

Some Formulae:

$$\text{Power} = \text{Capacity} \cdot \text{Voltage}^2 \cdot \text{Frequency}_{0 \rightarrow 1} + \text{Voltage} \cdot I_{\text{leakage}}$$

$$\text{Power} = \frac{\text{Joules}}{\text{Op}} \cdot \frac{\text{Ops}}{\text{Second}}$$

Dennard scaling: 0.7x voltage drop

Latency: How long it takes to do a task

Throughput: Total work done per unit of time e.g. queries/sec

Clock period: Duration of a clock cycle

Clock frequency: cycles per second

$$\text{Execution Time} = \text{Cycles per Program} \cdot \text{Clock Cycle Time} = \frac{\text{Cycles Per Program}}{\text{Clock Rate}}$$

$$\text{Clock Cycles} = \text{Instruction Count} \cdot \text{Cycles per Instruction}$$

$$\text{CPU Time} = \text{Instruction Count} \cdot \text{CPI} \cdot \text{Clock Cycle Time} = \frac{IC \cdot CPI}{\text{Clock Rate}}$$

$$\text{Exec Time} = \frac{\text{Instructions}}{\text{Program}} \cdot \frac{\text{Clock Cycles}}{\text{Instruction}} \cdot \frac{\text{Seconds}}{\text{Clock Cycle}}$$

$$\text{Clock Cycles} = \sum_{i=1}^n (CPI_i \cdot IC_i)$$

$$\text{Weighted average CPI} = \frac{\text{Clock Cycles}}{\text{Instruction Count}} = \sum_{i=1}^n \left(CPI_i \cdot \frac{IC_i}{IC} \right)$$

$$\text{Performance} = \frac{1}{\text{Exec Time}}$$

$$“X \text{ is } n \text{ faster than } Y” \text{ means } \frac{Perf_X}{Perf_Y} = \frac{Exec_Y}{Exec_X} = n$$

Energy = Average Power x Execution Time

$$\text{Optimize for Energy per Instruction: } Power = \frac{energy}{second} = \frac{energy}{instruction} \cdot \frac{instructions}{second}$$

$$\text{Amdahl's Law: } Speedup = \frac{CPUTime_{old}}{CPUTime_{new}} = \frac{CPUTime_{old}}{CPUTime_{old}[(1-f_x) + \frac{f_x}{S_x}]} = \frac{1}{(1-f_x) + \frac{f_x}{S_x}}$$

$$\text{Parallel Speedup: } Speedup = \frac{1}{(1-P) + \frac{P}{n}} \rightarrow P = \frac{\frac{Speedup}{1} - 1}{n - 1}$$

$$\text{Arithmetic Mean: } \frac{1}{n} \sum_{i=1}^n T_i \text{ used with times, not rates}$$

$$\text{Harmonic mean: } \frac{n}{\sum_{i=1}^n \frac{1}{R_i}} \text{ used with rates not times}$$

$$\text{Geometric mean: } \left(\prod_{i=1}^n \frac{T_i}{T_{ri}} \right)^{\frac{1}{n}} = \exp \left(\frac{1}{n} \sum_{i=1}^n \log \left(\frac{T_i}{T_{ri}} \right) \right)$$

$$\text{Power/performance benchmark: Overall ssj_ops per Watt} = \frac{\left(\sum_{i=0}^{10} ssj_ops_i \right)}{\left(\sum_{i=0}^{10} Power_i \right)}$$

Perf metrics:

time <application> measures execution time

perf record does low overhead sampling

perf topdown uses hardware performance counters

Flip flop setup time: $t_{ck} > t_{pd} + t_s + t_{skew}$

CISC:

Multi-cycle complex instructions

Load/store incorporated in instruction

Small code size

High CPI

Low clock frequency

Variable length instructions

RISC:

Simple (single-clock) Instructions

Register-to-register separate load instructions

Large code size

Low CPI

High clock frequency

Same length instructions

Simple instruction decode

Instructions:

R-type: [funct7] [rs2] [rs1] [funct3] [rd] [opcode]

I-type: [immediate[11:0]] [rs1] [funct3] [rd] [opcode]

S/B-type: [imm[11:5]] [rs2] [rs1] [funct3] [imm[4:0]] [opcode]

U/J-type: [immediate[31:12]] [rd] [opcode]

Call and Return:

Caller:

Save caller-saved registers as needed

Load arguments

Execute **JAL**

Callee setup:

Allocate memory for new frame (**xsp** = **xsp** - frame)

Save callee-saved registers as needed

Set frame pointer (**xfp** = **xsp** + frame size - 4)

Callee return:

Place return values in **x10** and **x11**

Restore any callee-saved registers

Pop stack (**xsp** = **xsp** + frame size)

Return by **jr xra**

Caller:

Restore any caller-saved registers as needed

Pipeline stages:

IF: Instruction fetch

ID: Instruction decode, register read

EX: execute operation or calculate address

MEM: Access memory command

WB: Write result back to register

Hazards:

Structure Hazards: A required hardware resource is busy

Data hazards: Must wait for previous instructions to produce/consume data

Read after Write: Instruction j tries to read before instruction i tries to write it

Write after Write: Instruction j tries to write an operand before i writes its value

Write after Read: Instruction j tries to write a destination before it is read by i

Control hazards: Next PC depends on current instruction result

Forwarding:

Identify *producers*: EX and MEM stages

All stages after first producer are sources of forwarding data: MEM and WB

Identify *consumers*: EX and MEM

These are the destinations of forwarded data

Forwarding and Hazard Detection:

```

if (MEM/WB.RegWrite && (MEM/WB.RegisterRd != 0) &&
    !(EX/MEM.RegWrite && (EX/MEM.RegisterRd != 0) &&
        (EX/MEM.RegisterRd = ID/EX.RegisterRs1)) && (MEM/WB.RegisterRd = ID/EX.RegisterRs1))
    ForwardA = 01
if (MEM/WB.RegWrite && (MEM/WB.RegisterRd != 0) &&
    !(EX/MEM.RegWrite &&
        (EX/MEM.RegisterRd != 0) && (EX/MEM.RegisterRd = ID/EX.RegisterRs2)) &&
        (MEM/WB.RegisterRd = ID/EX.RegisterRs2))
    ForwardB = 01
if (ID/EX.MemRead && ((ID/EX.RegisterRd = IF/ID.RegisterRs1) ||
    (ID/EX.RegisterRd = IF/ID.RegisterRs2)))
    stall the pipeline

```

Branch Prediction:

Static: predict not-taken

Dynamic:

Branch History Table: one entry for each branch, taken/not taken record

Branch Target Buffer: One entry for each branch, computes target address

Exceptions and Interrupts:

Switch from ‘user’ to ‘kernel’ mode

Exceptions: Arises within CPU. Undefined opcode, overflow, etc

Interrupt: External I/O controller, network card, etc

On exceptions:

Pass to relevant handler

Then return to program using EPC if possible

All previous instructions completed

Faulting instruction not started

No side effects

Nullifying instructions: Converts them to a NOP

Going past the 5-stage pipeline:

Instruction-level Parallelism:

Independence among instructions

Fetch multiple instructions per cycle

Evaluate which ones can go down the pipe together

Superscalar:

Double up on hardware in order to execute multiple instructions simultaneously

Deeper Pipelines: Increase the number of stages, decrease amount of logic. Higher clock freq.

Register Renaming: Map architectural registers to physical registers in decode stage to get rid of false dependencies. Need more physical registers than architectural ones.

Data flow graph:

Scoreboard: bit array, 1-bit for each GPR

If the bit is not set, the reg has valid data

If the bit is set, the reg has stale data

Dispatch in order: $RD \leftarrow F_n(RS, RT)$

If SB[RS] or SB[RT] is set \rightarrow RAW, stall

If SB[RD] is set \rightarrow WAW, stall

Else dispatch to functional unit, set SB[RD]

Complete out-of-order

Update GPR[RD], clear SB[RD]

Caching:

AMAT: Access Time = hit time + miss rate \cdot miss penalty

Three C’s of misses:

Compulsory: First time you’ve accessed this item

Capacity: Not enough room in the cache to hold item

Conflict: Item was replaced because of a conflict in its set

Direct-mapped cache:

2^n bytes total with 2^m byte blocks

Byte select: lower m bits

Cache index: lower $(n - m)$ bits of the memory address

Cache tag: upper $32 - n$ bits of the memory address

N -way set associative:

Each memory block can go to one of N entries in the cache

2^n -byte cache, 2^m -byte blocks, 2^a set-associative:

Cache contains $2^n / 2^m = 2^{n-m}$ blocks

Each cache way contains $2^{n-m} / 2^a = 2^{n-m-a}$ blocks

Byte offset: lowest m bits

Cache index: next $n - a$ bits

Associative caches might use Least Recently Used table for evictions

For N -way cache, $N!$ orderings

Write-through:

Main memory updated each cache write

Replacing a cache entry just overwrites new block

Memory write may cause pipeline stalls

Misses are simpler and cheaper

Uses a write buffer — FIFO queue

Write-back:

Only the cache entry is updated on each cache write

Cache and memory entries are inconsistent

Add ‘dirty’ bit to indicate whether memory needs to be updated

Write new value to memory on evictions

Writes are super fast

Write miss options: Do you allocate for space in the cache on a miss?

Do you fetch the rest of the block contents from memory?

For no-fetch-on-miss must use fine-grained valid bits

Write-back: typically write-allocate, fetch-on-miss

Multilevel caches:

Primary L1 caches attached to CPU: small, fast, focuses on hit time

Secondary L2 caches service misses from L1: larger, slower, still faster than DRAM

Cache Coherency:

P writes X , P reads $X \rightarrow$ read returns written value

P_1 writes X , P_2 reads X later \rightarrow read returns written value

P_1 writes X , P_2 writes $X \rightarrow$ all processors see writes in the same order

Single-Writer, Multiple-Read Invariant: For any memory location A , at any given epoch, there exists only one CPU that may write to A or some number of CPUs that may only read A

Data-Value Invariant: The value of the memory location at the start of an epoch is the same as the value of the memory location at the end of its last read-write epoch

CPU uses snooping protocols to ensure this:

2-state: very simple hardware and protocol. Write-through, all writes go on interconnect bus.

3-state MSI: Modified (one cache has valid copy), Shared (one or more have read-only copy),

Invalid (invalidated so one copy can go to modify state)