

Pete Wilcox  
CruzID: pcwilcox

## Tuples and Relations:

- *Tuple*: a  $k$ -tuple is an ordered sequence of  $k$  values
- If  $D_1, D_2, \dots, D_k$  are sets of elements then the cartesian product  $D_1 \times D_2 \times \dots \times D_k$  is the set of all  $k$ -tuples  $(d_1, d_2, \dots, d_k)$  such that  $\forall 1 \leq i \leq k: d_i \in D_i$
- *Relation*:
  - A  $k$ -ary relation is a subset of  $D_1 \times D_2 \times \dots \times D_k$  where each  $D_i$  is a set of elements
  - $D_i$  is the *domain* (or *datatype*) of the  $i^{th}$  column of the relation
  - Domains may be enumerated  $\{“AMS”, “CMPS”, “TIM”\}$  or may be of standard types
- An *attribute* is the name of a column in a relation
- A *relation schema*  $R$  is a set  $\{A_1, \dots, A_k\}$  of attributes written  $R(A_1, \dots, A_k)$ , where  $A_i$  is the name of the  $i^{th}$  column.
- A *relation database schema* or *database schema* is a set of relation schemas with disjoint relation names.

## SQL Primitives:

- **CHAR(n)**: fixed-length string of up to  $n$  characters (blank-padded with trailing spaces)
- **VARCHAR(n)**: also a string of up to  $n$  characters
- **BIT(n)**: padded on the right with 0s.
- **BIG VARYING(n)**: works like **VARCHAR**
- **BOOLEAN**: true, false, unknown
- **INT** or **INTEGER**: works like in C
- **SHORTINT**: works like **short int**
- **DECIMAL(n, d)**, **NUMERIC(n, d)**: total of  $n$  digits,  $d$  of them to the right of the decimal point
- **FLOAT(p)**, **FLOAT**, **REAL**
- **DOUBLE PRECISION**: analagous to **double** in c
- **DATE**, **TIME**, **TIMESTAMP**, **INTERVAL**: constants are character strings of specific form e.g. **DATE** ‘2017-09-13’
  - Subtracting one **TIME** from another results in an **INTERVAL**
  - Taking a **TIME** and adding an **INTERVAL** results in a **TIME**
  - Similarly for **TIMESTAMP** and **DATE**

## Tables:

- A *key constraint* or *key* of a relation schema  $R$  is a subset  $K$  of the attributes of  $R$  such that:
  1. For every instance  $r$  of  $R$ , every two distinct tuples of  $r$  must differ in their values of  $K \iff$  there can’t be two different tuples that have the same value for key  $K$
  2. No proper subset of  $K$  has the above property
- A *superkey* is a set of attributes of  $R$  that includes a key of  $R$
- **CREATE TABLE R(A, B, C, PRIMARY KEY(A))**:
  1. None of the tuples in  $R$  can have null  $A$  values
  2. Rows are uniquely identified by their  $A$  values
  3. There can be at most one primary key for a table

- **CREATE TABLE S(D, E, F, UNIQUE(D))**:
  1. Rows in  $S$  can contain null  $D$  values
  2. Rows with *non-null*  $D$  values are uniquely identified by their  $D$  values
  3. There can be multiple unique constraints in addition to a primary key
- **CREATE TABLE T(G NOT NULL, H DEFAULT ‘foo’)**:
  1. If no default value is specified and no value is entered then the value will be **NULL**
  2. **NOT NULL** prevents a column from having null values
  3. If a default value is specified and no value is entered then the value will be the default

## Queries:

- Basic form:  
**SELECT [DISTINCT] c1, c2, ..., cm**  
**FROM R1, R2, ..., Rn**  
**[WHERE condition]**
- **SELECT**:
  - Projection: **SELECT title, year** - only a subset of attributes from the relation(s) in the **FROM** clause is selected
  - **DISTINCT**: Removes duplicate tuples from result
  - Aliasing: **SELECT title AS name** - rename the attributes in the result
  - Expressions are allowed in the **SELECT** clause. Ex: **SELECT title AS name, length \* 60 AS durationInSeconds**
  - Constants can also be included: **SELECT title AS name, length \* 60 AS durationInSeconds, ‘seconds’ AS inSeconds**
- **WHERE**:
  - Comparison operators: **=, <>, <, >, <=, >=**
  - Logical connectives: **AND, OR, NOT**
  - Arithmetic expressions: **+, -, \*, /**, etc
  - In general the **WHERE** clause is a boolean expression where each condition is of the form *expression op expression*
- Pattern matching with the **LIKE** operator:
  - **s LIKE p, s NOT LIKE p**
  - **s** is a string, **p** is a pattern
  - **‘\%’** stands for 0 or more arbitrary characters
  - **‘\_’** stands for exactly one arbitrary character
  - Matching quotes: **WHERE x LIKE ‘’’’** matches one
  - Matching quotes: **WHERE x LIKE ‘’’’** matches two
  - Matching **\%** or **\_**: **WHERE x LIKE ‘!\%!\%!’ESCAPE ‘!’** where **!** can be any character
- **DATE** and **TIME** and **TIMESTAMP**
  - Separate data types
  - Constants are character strings of the form:  
**DATE** ‘2015-01-13’  
**TIME** ‘16:45:33’  
**TIMESTAMP** ‘2015-01-13 16:45:33’
  - **DATE, TIME, TIMESTAMP** can be compared using ordinary

comparison operators e.g. **WHERE ReleaseDate <= DATE ‘1990-06-19’**

- If **Salary** is **NULL** then the following will be **UNKNOWN**:
  - **Salary = 10**
  - **Salary <> 10**
  - **90 > Salary OR 90 <= Salary**
  - **Salary = NULL**
  - **Salary <> NULL**
- Use of **IS NULL** and **IS NOT NULL**:
  - **Salary IS NULL** will be true if **SALARY is NULL**, false otherwise
  - **Salary IS NOT NULL** will be true if **SALARY is not NULL**, false otherwise
- Ordering the result:
  - **ORDER BY** presents the result in a sorted order
  - By default the result will be ordered in ascending order **ASC**
  - For descending order on an attribute you write **DESC** in the list of attributes
- Multiple relations in **FROM** clause: for every tuple  $t_1 \in R_1, t_2 \in R_2, \dots, t_n$  from  $R_n$  if  $t_1, \dots, t_n$  satisfy *condition* then add the resulting tuple that consists of  $c_1, c_2, \dots, c_m$  components of  $t$  into the result

## Joins: With relations R(A,B,C) and S(C,D,E)

- **R JOIN S ON R.B=S.D AND R.A=S.E**:
  - Selects only tuples from  $R$  and  $S$  where  $R.B=S.D$  and  $R.A=S.E$
  - Schema of the resulting relation: **(R.A, R.B, R.C, S.C, S.D, S.E)**
  - Equivalent to:  
**SELECT \***  
**FROM R, S**  
**WHERE R.B=S.D AND R.A=S.E;**
- **R CROSS JOIN S**:
  - Product of the two relations  $R$  and  $S$
  - Schema of the resulting relation: **(R.A, R.B, R.C, S.C, S.D, S.E)**
  - Equivalent to:  
**SELECT \***  
**FROM R, S;**
- **R NATURAL JOIN S**:
  - Schema of the resulting relation: **(A, B, C, D, E)**
  - Equivalent to:  
**SELECT R.A, R.B, R.C, S.D, S.E**  
**FROM R, S**  
**WHERE R.C = S.C**

## Set and Bag Operations: R(A,B,C), S(A,B,C)

- **UNION**: Set union
  - Input to union must be *union-compatible*:  $R$  and  $S$  must have attributes of the same type, in the same order
  - Output of the union has the same schema as  $R$  or  $S$

- Meaning: Output consists of the *set* of all tuples from  $R$  and from  $S$
- Could (should?) have been called **UNION DISTINCT**  
(**SELECT \* FROM R**)  
**UNION**  
(**SELECT \* FROM S**)
- **UNION ALL**: Bag union
  - Input must be *union-compatible*
  - Output has the same schema as  $R$  or  $S$
  - Output consists of the collection of all tuples from  $R$  and from  $S$  *including duplicates*.
  - Attributes/column names may be different -  $R$ 's are used
- **INTERSECT, INTERSECT ALL**: set/bag intersection
  - Input must be union-compatible.
  - $Query_1$  **INTERSECT**  $Query_2$
  - $Query_1$  **INTERSECT ALL**  $Query_2$
  - Find all tuples that are in the results of both  $Query_1$  and  $Query_2$ .
  - **INTERSECT** is distinct. **INTERSECT ALL** reports duplicates.
- **EXCEPT, EXCEPT ALL**: set difference, bag difference
  - Must be union-compatible
  - $Query_1$  **EXCEPT**  $Query_2$
  - $Query_1$  **EXCEPT ALL**  $Query_2$
  - Find all tuples that are in the result of  $Query_1$  and not in the result of  $Query_2$
  - **EXCEPT** is distinct, **EXCEPT ALL** is not
- Order of operations: **INTERSECT** has higher precedence than **UNION** and **EXCEPT**.

### Subqueries:

- A query embedded in another query
- Can be used as a boolean or can return a constant or can return a relation
- **IN, NOT IN**: used to select from subquery that returns relation
- **WHERE A < ANY**: checks that attribute  $A$  is less than at least one of the answers returned by the subquery.
- **EXISTS**: Checks that subquery returns non-empty result. Also: **NOT EXISTS**

### Aggregates and Grouping:

- Basic SQL has 5 aggregation operators: **SUM, AVG, MIN, MAX, COUNT**
- Aggregation operators work on scalar values, except for **COUNT(\*)** which counts the number of tuples
- **GROUP BY** clause follows the **WHERE** clause
  - Let Result begin as an empty multiset of tuples
  - For every tuple  $t_1$  from  $R_1$ ,  $t_2$  from  $R_2, \dots, t_n$  from  $R_n$ : if  $t_1, \dots, t_n$  satisfy *condition* then add the resulting tuple that consist of  $c_1, c_2, \dots, c_m$  of the  $t_i$  into Result
  - Group the tuples according to the grouping attributes - if **GROUP BY** is omitted, the entire table is one group
- **NULLs** are ignored in any aggregation
  - They do not contribute to the **SUM, AVG, COUNT, MIN, MAX** of an attribute
  - **COUNT(\*)** = the number of tuples in a relation even if some columns are **NULL**
  - **COUNT(A)** is the number of tuples with *non-NULL* values for  $A$
  - **SUM, AVG, MIN, MAX** on an empty result (no tuples) is **NULL**
  - **COUNT** of an empty result is 0
  - **GROUP BY** does *not* ignore **NULL**
- **HAVING** clause:
  - Choose groups based on some aggregate property of the group itself
  - Same attributes and aggregates that can appear in the **SELECT** can appear in the **HAVING** clause condition
  - Can use **EVERY** to constrain **HAVING** to all tuples in the group e.g. **HAVING COUNT(\*) > 1 AND EVERY (S.age <= 40)**

### Database Modification Statements:

- **INSERT INTO R(A1, ..., An)VALUES (v1, ..., vn)**: a tuple  $(v_1, \dots, v_n)$  is inserted into  $R$  such that  $A_i = v_i \forall i$  and default values (perhaps **NULL**) are entered for any missing attributes.
- **DELETE FROM R WHERE <condition>**: Deletes *all* tuples such that the condition evaluates as true - if there is no **WHERE** clause it will delete all tuples in  $R$
- **UPDATE R SET <new-value-assignments> WHERE <condition>**: Change the given attribute to the new value

in every tuple in  $R$  where the condition is true

- Semantics: database modifications are completely evaluated on the old state of the database producing a new state of the database

### Transaction:

- Transactions provide ACID properties: atomicity, consistency, isolation, durability
- **START TRANSACTION** or **BEGIN TRANSACTION**: marks the beginning of a transaction, followed by one or more SQL statements
- **COMMIT**: Ends the transaction. All changes are durably written to the backing store and become visible to other transactions.
- **ROLLBACK**: Causes the transaction to abort or terminate. None of the changes are committed.
- **SET TRANSACTION READ ONLY**:
  - set *before* the transaction begins, tells the SQL system that the next transaction is read-only
  - SQL uses this to parallelize many read-only transactions
- **SET TRANSACTION READ WRITE**:
  - Tells SQL that the next transaction may write data in addition to read
  - Default option if not specified, often not specified
- Dirty Reads: *Dirty data* refers to data that is written by a transaction but has not yet been committed by the transaction
- Isolation levels:
  - **SET TRANSACTION READ WRITE ISOLATION LEVEL READ UNCOMMITTED**
  - Default isolation level depends on system, most run with **READ COMMITTED** or **SNAPSHOT ISOLATION**
  - **READ COMMITTED**: only clean(committed) reads but you might read data committed by other transactions
  - **REPEATABLE READ**: repeated queries of a tuple during a transaction will retrieve the same value. Also, a second scan may return 'phantoms' which are tuples newly inserted while the transaction is running.
  - **SERIALIZABLE**: Can be replayed one by one.