

ACID: atomicity, consistency, isolation, durability

Network Models: Network, Hierarchical, Relational

Tuples and Relations:

- *Tuple*: a k -tuple is an ordered sequence of k values
- If D_1, D_2, \dots, D_k are sets of elements then the cartesian product $D_1 \times D_2 \times \dots \times D_k$ is the set of all k -tuples (d_1, d_2, \dots, d_k) such that $\forall 1 \leq i \leq k: d_i \in D_i$
- *Relation*:
 - A k -ary relation is a subset of $D_1 \times D_2 \times \dots \times D_k$ where each D_i is a set of elements
 - D_i is the *domain* (or *datatype*) of the i^{th} column of the relation
 - Domains may be enumerated $\{“AMS”, “CMPS”, “TIM”\}$ or may be of standard types
- An *attribute* is the name of a column in a relation
- A *relation schema* R is a set $\{A_1, \dots, A_k\}$ of attributes written $R(A_1, \dots, A_k)$, where A_i is the name of the i^{th} column.
- A *relation database schema* or *database schema* is a set of relation schemas with disjoint relation names.

Declarative vs Procedural: A *declarative* language places constraints on the output but not on *how* that output is obtained. A *procedural* language specifies a sequence of operations to obtain the desired output. SQL is not fully declarative.

SQL DDL and SQL DML:

- *Data Definition Language*: CREATE TABLE, DROP TABLE, CREATE SCHEMA, DROP SCHEMA
- *Data Manipulation Language*: SELECT, INSERT, UPDATE, DELETE

Primitives:

- CHAR(n): fixed-length string of up to n characters (blank-padded with trailing spaces)
- VARCHAR(n): also a string of up to n characters
- BIT(n): padded on the right with 0s.
- BIG VARYING(n): works like VARCHAR
- BOOLEAN: true, false, unknown
- INT or INTEGER: works like in C
- SHORTINT: works like `short int`
- DECIMAL(n , d), NUMERIC(n , d): total of n digits, d of them to the right of the decimal point
- FLOAT(p), FLOAT, REAL
- DOUBLE PRECISION: analagous to `double` in c
- DATE, TIME, TIMESTAMP, INTERVAL: constants are character strings of specific form e.g. DATE ‘2017-09-13’
 - Subtracting one TIME from another results in an INTERVAL
 - Taking a TIME and adding an INTERVAL results in a TIME
 - Similarly for TIMESTAMP and DATE

Tables:

- A *key constraint* or *key* of a relation schema R is a subset K of the attributes of R such that:
 1. For every instance r of R , every two distinct tuples of r must differ in their values of $K \iff$ there can’t be two different tuples that have the same value for key K
 2. No proper subset of K has the above property
- A *superkey* is a set of attributes of R that includes a key of R
- CREATE TABLE R(A, B, C, PRIMARY KEY(A)):
 1. None of the tuples in R can have null A values
 2. Rows are uniquely identified by their A values
 3. There can be at most one primary key for a table
- CREATE TABLE S(D, E, F, UNIQUE(D)):
 1. Rows in S can contain null D values
 2. Rows with *non-null* D values are uniquely identified by their D values
 3. There can be multiple unique constraints in addition to a primary key
- CREATE TABLE T(G NOT NULL, H DEFAULT ‘foo’):
 1. If no default value is specified and no value is entered then the value will be NULL
 2. NOT NULL prevents a column from having null values
 3. If a default value is specified and no value is entered then the value will be the default

Queries:

- Basic form:
SELECT [DISTINCT] c_1, c_2, \dots, c_m
FROM R_1, R_2, \dots, R_n
[WHERE condition]
- SELECT:
 - Projection: SELECT title, year - only a subset of attributes from the relation(s) in the FROM clause is selected
 - DISTINCT: Removes duplicate tuples from result
 - Aliasing: SELECT title AS name - rename the attributes in the result
 - Expressions are allowed in the SELECT clause. Ex: SELECT title AS name, length * 60 AS durationInSeconds
 - Constants can also be included: SELECT title AS name, length * 60 AS durationInSeconds, ‘seconds’ AS inSeconds
- WHERE:
 - Comparison operators: =, <>, <, >, <=, >=
 - Logical connectives: AND, OR, NOT
 - Arithmetic expressions: +, -, *, /, etc
 - In general the WHERE clause is a boolean expression where each condition is of the form *expression op expression*
- Pattern matching with the LIKE operator:
 - s LIKE p , s NOT LIKE p
 - s is a string, p is a pattern
 - ‘\%’ stands for 0 or more arbitrary characters
 - ‘_’ stands for exactly one arbitrary character

- Matching quotes: WHERE x LIKE ‘’’’ matches one quote symbol
- Matching quotes: WHERE x LIKE ‘’’’’’ matches two quote symbols
- Matching \% or _: WHERE x LIKE ‘!\%!\’ESCAPE ‘!’ where ! can be any character
- DATE and TIME and TIMESTAMP
 - Separate data types
 - Constants are character strings of the form:
DATE ‘2015-01-13’
TIME ‘16:45:33’
TIMESTAMP ‘2015-01-13 16:45:33’
 - DATE, TIME, TIMESTAMP can be compared using ordinary comparison operators e.g. WHERE ReleaseDate <= DATE ‘1990-06-19’
- If Salary is NULL then the following will be UNKNOWN:
 - Salary = 10
 - Salary <> 10
 - 90 > Salary OR 90 <= Salary
 - Salary = NULL
 - Salary <> NULL
- Use of IS NULL and IS NOT NULL:
 - Salary IS NULL will be true if SALARY is NULL, false otherwise
 - Salary IS NOT NULL will be true if SALARY is not NULL, false otherwise
- Ordering the result:
 - ORDER BY presents the result in a sorted order
 - By default the result will be ordered in ascending order ASC
 - For descending order on an attribute you write DESC in the list of attributes
- Multiple relations in FROM clause: for every tuple $t_1 \in R_1, t_2 \in R_2, \dots, t_n \in R_n$ if t_1, \dots, t_n satisfy *condition* then add the resulting tuple that consists of c_1, c_2, \dots, c_m components of t into the result

JOINS : With relations R(A,B,C) and S(C,D,E)

- R JOIN S ON R.B=S.D AND R.A=S.E:
 - Selects only tuples from R and S where R.B=S.D and R.A=S.E
 - Schema of the resulting relation: (R.A, R.B, R.C, S.C, S.D, S.E)
 - Equivalent to:
SELECT *
FROM R, S
WHERE R.B=S.D AND R.A=S.E;
- R CROSS JOIN S:
 - Product of the two relations R and S
 - Schema of the resulting relation: (R.A, R.B, R.C, S.C, S.D, S.E)
 - Equivalent to:

p	q	p OR q	p AND q	p = q
T	T	T	T	T
T	F	T	F	F
T	U	T	U	U
F	T	T	F	F
F	F	F	F	T
F	U	U	F	U
U	T	T	U	U
U	F	U	F	U
U	U	U	U	U

- ```

SELECT *
FROM R, S;

```
- **R NATURAL JOIN S:**
    - Schema of the resulting relation: (A, B, C, D, E)
    - Equivalent to:

```

SELECT R.A, R.B, R.C, S.D, S.E
FROM R, S
WHERE R.C = S.C

```

### Set and Bag Operations: R(A,B,C), S(A,B,C)

- **UNION:** Set union
  - Input to union must be *union-compatible*: R and S must have attributes of the same type, in the same order
  - Output of the union has the same schema as R or S
  - Meaning: Output consists of the *set* of all tuples from R and from S
  - Could (should?) have been called **UNION DISTINCT**

```

(SELECT *
FROM R)
UNION

```

```

(SELECT *
FROM S)

```

- **UNION ALL:** Bag union
  - Input must be *union-compatible*
  - Output has the same schema as R or S
  - Output consists of the collection of all tuples from R and from S *including duplicates*.
  - Attributes/column names may be different - R's are used
- **INTERSECT, INTERSECT ALL:** set intersection, bag intersection
  - Input must be union-compatible.
  - *Query<sub>1</sub>* **INTERSECT** *Query<sub>2</sub>*
  - *Query<sub>1</sub>* **INTERSECT ALL** *Query<sub>2</sub>*
  - Find all tuples that are in the results of both *Query<sub>1</sub>* and *Query<sub>2</sub>*.
  - **INTERSECT** is distinct. **INTERSECT ALL** reports duplicates.
- **EXCEPT, EXCEPT ALL:** set difference, bag difference

- Must be union-compatible
- *Query<sub>1</sub>* **EXCEPT** *Query<sub>2</sub>*
- *Query<sub>1</sub>* **EXCEPT ALL** *Query<sub>2</sub>*
- Find all tuples that are in the result of *Query<sub>1</sub>* and not in the result of *Query<sub>2</sub>*
- **EXCEPT** is distinct, **EXCEPT ALL** is not
- Order of operations: **INTERSECT** has higher precedence than **UNION** and **EXCEPT**.

### Subqueries:

- A query embedded in another query
- Can be used as a boolean or can return a constant or can return a relation
- **IN, NOT IN:** used to select from subquery that returns relation
- **WHERE A < ANY:** checks that attribute A is less than at least one of the answers returned by the subquery.
- **EXISTS:** Checks that subquery returns non-empty result. Also: **NOT EXISTS**