

- **ACID**
 - **ATOMICITY**: An atomic transaction happens as one unit, either the whole thing commits or none of it does.
 - **CONSISTENCY**: A consistent transaction brings the DB from one valid state to another valid state with respect to any constraints.
 - **ISOLATION**: Concurrent isolated transactions would have the same result if run sequentially.
 - **DURABILITY**: A committed transaction will remain committed even in the event of a hardware failure.
- **RAID Levels**
 - Level 0: No redundancy (just stripin)
 - Level 1: Mirrored (two identical copies)
 - * Each disk has an exact mirror image
 - * Parallel reads; writes involve two disks
 - * Maximum transfer rate = transfer rate of one disk
 - Level 0+1 (Level 10): Striping and Mirroring
 - * Parallel reads; writes involve two disks
 - * Maximum transfer rate = aggregate bandwidth
 - Level 3: Bit-interleaved parity
 - * Striping Unit: one bit (or byte) (one check disk)
 - * Each read and write request involves all disks; disk array can process one request at a time
 - Level 4: Block-interleaved parity
 - * Striping unit: one disk block (one check disk)
 - * Parallel reads possible for small requests, large requests can utilize full bandwidth
 - * Writes involve modified block *and* check disk
 - Level 5: Block-interleaved distributed parity
 - * Similar to RAID level 4 but parity blocks are distributed over all disks
- **Buffer Management in a DBMS**
 - DBMS maintains buffer pool of frames, each frame holds a page, info is in **<frame#, pageid>** table
 - Choice of frame replacement dictated by replacement policy such as LRU
 - When a page is requested:
 - * If requested page is not in pool:
 - Choose a frame for replacement
 - If that frame is dirty, write it to disk
 - Read requested page into chosen frame
 - * Pin the page and return its address
 - * When done the requestor must indicate whether the page has been modified (dirty bit) and unpin
 - * Page in pool may be requested many times
 - A pin count is used and a page is a candidate for replacement iff **pin_count = 0**
 - Pinning increments pin count and unpinning decrements
 - * Concurrency control and recovery may entail additional I/O when a frame is chosen for replacement (write-ahead log protocol)
 - * Frame is chosen for replacement using LRU, clock, MRU, etc
 - * Sequential flooding: Caused by using LRU when the number of buffer frames is less than the number of pages in the file
- **Files of Records**
 - Page or block is ok when doing I/O but higher levels of DBMS operate on *records* and thus want *files of records*
 - **FILE**: A collection of pages each containing a collection of records. Must support
 - * Insert (append)/delete/modify record
 - * Read a particular record specified using *record id*
 - * Scan all records possibly with some conditions on the records to be retrieved
 - **Unordered ‘Heap’ Files**:
 - * Simplest file structure that contains records in no particular (logical) order
 - * As file grows and shrinks, disk pages are allocated and de-allocated
 - * To support record-level operations we must:
 - Keep track of the *pages* in a file: **page id (pid)**
 - Keep track of the *free space* on a page
 - Keep track of the *records* on a page: **record id (rid)**
 - Keep track of *fields* within records
 - * Operations: create/destroy file, insert/delete record, fetch record with specific **rid**, scan all records
 - Record formats: **Fixed Length**
 - * Information about field types is the same for all records in file; it is stored in *system catalogs*
 - * Finding the i^{th} field of a record does not require scanning the record
 - Record formats: **Variable length**
 - * Several alternative formats (# of fields is fixed)
 - * Fields delimited by special symbols (e.g. \$ between fields)
 - * Fields preceded by lengths
 - Record formats: **Variable length with directory**
 - * Use array of offsets at start of record
 - **Heap file implemented as a list**
 - * The header page id and heap file name must be stored someplace
 - * Each page contains two extra pointers in this case
 - * Refinement: use several lists for different degrees of free space
 - Page formats:
 - * File → collection of pages
 - * Page $-i$: collection of tuples/records
 - * Query operators deal with tuples
 - * Slotted page format:
 - Each page has a collection of *slots*
 - Each slot contains a record
 - * RID: **<page id, slot number>**
 - **Heap file using a page directory**
 - * Page entries can include the number of free bytes on each page
 - * Directory is a collection of pages; linked list is one possible implementation
 - **System catalogs**:
 - * For each relation:
 - name, file, file structure
 - name, type, length (if fixed) for each attribute
 - Index name, target, and kind for each index
 - also integrity constraints, defaults, nullability, etc
 - * For each index: structure (e.g. B+ tree) and search key fields
 - * For each view: view name and definition (including query)
 - * Plus statistics, authorization, buffer pool size, etc
 - **Column Stores**:
 - * Store data “vertically”
 - * Contrast with a “row-store” that stores all the attributes of a tuple/record contiguously
 - * Each column can be stored as a separate file and compressed
 - * SAP HANA:
 - Dictionary compression per column
 - Column main: read-optimized store for immutable data. Uses high data compression and heuristic algorithms to order data to maximize secondary compression
 - Column delta: write-optimized store for inserts, updates, deletes. Uses less compression, appends updates to the end, and merges with main periodically.
 - * Additional types: prefix coding, run length coding, cluster coding, sparse coding, indirect coding
 - **Indexes**:
 - Speeds up selections on the search key fields for the index
 - Contains a collection of data entries and supports efficient retrieval of all data entries k^* with a given key value k
 - **B+ Tree Indexes**
 - Leaf pages contain *data entries* and are chained (prev & next)
 - Non-leaf pages have *index entries*, used to direct searches
 - Insert/delete at $\log_F N$, keep tree *height-balanced* (F = fanout, N = # leaf pages)
 - Minimum 50% occupancy (in all nodes except root). Each node contains $d \leq m \leq 2d$ entries; d = the *order* of the tree.
 - Typical order $d = 100$
 - Percentage of node that is full is more useful, typical fill-factor 67%
 - Average *fanout* for non-leaves $F = 133$
 - Inserting a data entry:
 - * Find correct leaf L
 - * Put data entry onto L
 - * If L has enough space, done
 - * Otherwise, must split L . Redistribute entries evenly,

- copy up the middle key (key must still exist in leaf). Insert index entry pointing to L_2 into parent of L .
- * This can happen recursively: if parent of L grows, need to push up middle key.
- * Splits “grow” the tree; root split increases height.
- Deleting a data entry:
 - * Start at root, find leaf L where entry belongs
 - * Remove the entry
 - * If L is at least half full, done
 - * Otherwise, if L has only $d - 1$ entries, try to redistribute, borrowing from sibling (adjacent node with same parent)
 - * If redistribution fails, merge L and sibling
 - * If merge occurred, must delete entry from parent (pointing to merged node)
 - * Merge can propagate to root, decreasing height of the tree
- **Hash-Based Indexes**:
 - Good for equality selections
 - Index is a collection of *buckets*. Each bucket = *primary page* plus zero or more *overflow pages* (called *static hashing*). Buckets contain data entries.
 - *Hashing function h* : $h(r)$ = bucket in which (data entry for) record r belongs. h looks at the *search key* fields of r .
- Alternatives for Data Entry k^* in index:
 - In a data entry k^* we can store: an actual data record, or **<k, RID>**, or **<k, list of RIDs>**
 - Choice of alternative for entries is orthogonal to the indexing technique
- Alternative 1: data records live in index
 - Index structure is actually a file organization for the data records
 - At most one index on a given collection of data can use this Alternative
 - If data records are very large, # of leaf pages containing data entries is high.
- Alternatives 2 and 3: Key/RID or Key/RIDlist:
 - Data entries are typically much smaller than data records
 - Alternative 3 is more compact but leads to variable-sized data entries, even if the search keys are of fixed length
- Index classification:
 - *Primary vs Secondary*: if search key contains the primary key, index is called the primary index
 - *Clustered vs Unclustered*: If order of data records is the same as (or close to) the order of stored data records then index is called a clustered index.
- A back of the envelope cost model:
 - B : the number of data pages
 - R : number of records per page
 - D : average time to read or write a disk page
 - F : average fanout for a non-leaf page
- Indexes with composite search keys:
 - Composite search keys: search on a combination of fields
 - Equality query: every field value is equal to a constant value
 - Range query: some field value doesn’t have equality test
 - Data entries in index sorted by search key to support range queries
- **ISAM**: Index-Sequential Access Method
 - Index file has first key on each page, can binary search index then scan the page.
 - *Static* structure, inserts and deletes only affect leaf or overflow pages.
 - If index is very large, recursively create a second layer (and so on).
 - *File Creation*: Leaf pages first allocated sequentially, sorted by search key; then index pages allocated, and then overflow pages.
 - *Index entries*: **<key value, page id>**; they ‘direct’ searches for *data entries* which are in leaf pages
 - *Search*: Start at root; use key comparisons to go to leaf. I/O cost $\propto \log_F N$ where F = # entries/index pg, N = # leaf pgs
 - *Insert*: Find leaf where data entry belongs and put it there, using overflow page if necessary.
 - *Delete*: Find and remove from leaf; if empty overflow
- page, deallocate
- **B-Tree Prefix Key Compression**: Increase fan-out by reducing the size of search keys on interior nodes. key values only direct traffic so we only need the minimum length for that
- **Bulk Loading of a B+ Tree**
 - Creating a new B+ tree by inserting one at a time is very slow, bulk loading is better
 - *Initialization*: Sort all data entries, insert pointer to first (leaf) page in a new (root) page
 - Index entries for leaf pages always entered into right-most index page just above leaf level. When this fills up it splits.
- **Log-Structured Merge Tree**: Sequential trees of exponentially larger size. Inserts go to smallest smallest tree, deletes insert tombstone records, spill to next-deeper level on overflow
- **R-Tree**: Tree of rectangles, search for intersections between them
- **Static Hash-based Indexes**:
 - # primary pages is fixed, allocated sequentially, never de-allocated; overflow pages if needed
 - $h(k) \bmod M$ = bucket(page) to which data entry with key k belongs (M = # buckets)
 - Buckets contain data entries
 - Hash function works on *search key field* of record r . Must distribute over range $0 \dots M - 1$
 - $h(key) = (a * key + b)$ usually works well; a and b are constants to tune h
 - *Long overflow chains* can develop and degrade performance.
- **Extendible Hashing**:
 - Situation: bucket (primary page) becomes full. Solved by doubling number of buckets instead of using an overflow page
 - Use directory with pointers to buckets. Double the number of buckets by doubling the directory and splitting buckets as needed
 - Only one bucket at a time splits. No overflow pages
 - *Global Depth* is the last d bits after hashing and indexes into the directory to determine which bucket is used
 - *Local Depth* is used for each bucket. If the $LD == GD$ and the bucket splits, the directory must double.
- **Bitmap Indexes**:
 - Index which allows for fast equality checks. Order the records in some $O(1)$ way and maintain one or more bit vectors storing their values for particular fields.
 - One bitmap for each distinct domain value, and one bitmap for NULL if the column can be null.
 - Can use XOR operation to reduce number of maps needed by one.
- **External Sorting**:
 - Goal: Need to sort more data than will fit in memory, efficiently.
 - **2-Way Sort**: Requires 3 buffers
 - * Pass 0: read a page, sort it, write it out (only one buffer page used)
 - * Pass 1, 2, 3, ...: Read and merge pairs of runs. (Three buffer pages are used)
 - Sorting $N = 2^k$ *Pages of Data*:
 - * Pass 0: read, sort, write $\rightarrow 2^k$ 1-page runs
 - * Pass 1: Read + merge 1-page pairs, write $\rightarrow 2^{k-1}$ 2-page runs
 - * Pass 2: Read + merge 2-page pairs, write $\rightarrow 2^{k-2}$ 4-page runs
 - * Pass $k - 1$: Read + merge 2^{k-2} -page pairs, write $\rightarrow 2^{k-1}$ -page runs
 - * Pass k : Read + merge 2^{k-1} -page pairs, write $\rightarrow 12^k$ -page result
 - 2-Way External merge sort: N pages in file $\implies \lceil \log_2 N \rceil + 1$ passes, total I/O cost = $2N (\lceil \log_2 N \rceil + 1)$
 - *General external merge sort*:
 - * Sorting a file with N pages using B buffer pages.
 - * Pass 0: use B buffer pages. Produce $\lceil N/B \rceil$ sorted runs of B pages each.
 - * Pass 2, ... etc: merge $b - 1$ runs
 - * Number of passes: $1 + \lceil \log_B \lceil N/B \rceil \rceil$
 - * Cost = $2N * (\# \text{ of passes })$
 - *Double Buffering*

- * To reduce wait time for I/O request to complete, can *prefetch* into **shadow block**
- * Potentially more passes; in practice, most files still sorted in 2-3 passes.
- *B+ Tree as “Sorted Access Path”*
- * Scenario: table to be retrieved in some order has a B+ tree index on the ordering columns
- * Idea: retrieve records in order by traversing the B+ tree's leaf pages
- * Very good idea if the tree is clustered, otherwise probably a bad idea
- **Query Processing**
- *Access Paths:*
- * An access path is a method of retrieving tuples: file scan or index that matches a selection in the query
- * A tree index *matches* (a conjunction of) terms that involve only attributes in a *prefix* of the search key.
- * A hash index *matches* (a conjunction of) terms that has a term *attribute = value* for every attribute in the search key of the index.
- Selection conditions often first converted to be in CNF (ANDing of ORs)
- One approach to selections:
- * Find the *most selective access path*, retrieve tuples using it, then apply any remaining terms which don't match the index
- * Most selective access path: an index or file scan that we estimate will require the fewest page I/Os
- * Terms that match this index reduce the number of tuples retrieved; other terms used to filter the retrieved tuples on the fly, but don't prevent retrieval of the tuples/pages
- Using an index for selections: cost depends on # qualifying tuples and clustering
- * Cost of finding qualifying data entries (typically small) plus cost of retrieving the actual records themselves (can be large without clustering)
- *Duplicate Elimination*
- * Relational algebra projection removes duplicates: SQL systems don't remove duplicates unless the keyword **DISTINCT** is specified
- * Sorting approach: sort on <sid, bid> and remove duplicates. Can optimize by dropping unwanted columns while sorting.
- * Hashing approach: hash on <sid, bid> to create *partitions*. Load partitions into memory one at a time, build an in-memory hash structure, and eliminate duplicates within it.
- Notation:
- * Pages in a heap relation R : Pages_R
- * Tuples per page for a relation R : TPP_R
- * Number of tuples in R : Card_R
- * $\text{Card}_R = \text{Pages}_R * \text{TPP}_R$
- *Simple Nested Loops Join*: foreach tuple in R , for each tuple in S , if $r_i == s_j$ then add <**r,s**> to result.
- * For each tuple in the *outer* relation R we scan the entire *inner* relation S . Cost: $\text{Pages}_R + \text{Card}_R * \text{Pages}_S$
- * Page-oriented nested loops join: for each *page* of R get each *page* of S and write out matching pairs of tuples
- *Index Nested Loops*
- * If there is an index on the join column of one relation,

- can make it the inner and exploit the index. Cost: $\text{Pages}_R + \text{Card}_R$ * cost of finding matching S tuples
- * For each R tuple, cost of probing S index is about 1.2 for hash index, 2-4 for B+ tree. Cost of then finding S tuples (assuming alt. 2 or 3) depends on clustering. Clustered typically 1 I/O, unclustered up to 1 I/O per matching S tuple.
- *Block Nested Loops Join*
- * Use one page as an input buffer for scanning the inner S , one pas as the output buffer, and use *all* remaining pages to hold “block” of pages of outer R .
- * For each block of R , hash each data entry to a hash table. Then compare all entries in S .
- *Join: Sort-Merge ($R \bowtie_{i=j} S$)*
- * Sort R and S on the join column, then scan them to do a ‘merge’ (on join column) and then output result tuples.
- * R is scanned once; each S group is scanned once per matching R tuple.
- *Statistics and Catalogs*
- * Catalogs typically contain at least: # *tuples* and # *pages* in each relation; # *distinct key values* and # *pages* for each index; *index height*, *low/high key values* for each tree index
- * Catalogs updated periodically: updating each time data changes is too expensive and approximation is fine
- *Cost estimation*: For each plan must estimate cost
- * Estimate *Cost* of each operation in plan-tree: depends on input cardinalities
- * Estimate *Size* of result for each operation in the tree; for seelctions and joins assume independence of predicates
- *Size estimation and reduction factors*
- * Maximum # of tuples in result is the product of cardinalities in the **FROM** clause
- * Reduction Factor RF associated with each *term* reflects the impact of *term* in reducing result size
- *Grace Hash-Join*
- * Like a two-phase index nested loop join
- * *Build phase*: Partition both relations using hash function h : R tuples in partition i will only match S tuples in partition i .
- * *Match phase*: Read in a partition of R , hash it using $h_2(<> h)$. Scan matching partition of S searching for its R matches.
- * # partitions $k \leq B - 1$ and Pages_R/k (size of largest partition to be held in memory) $< B - 1$.
- * Assuming uniform-sized partitions and maximizing k : $k = B - 1$ and $\text{Pages}_R/k < B - 1$ so $(B - 1)^2 > \text{Pages}_R \implies B > \sqrt{\text{Pages}_R}$
- * Can hash-join *recursively* to reduce the amount of memory needed
- * In build phase: Read + write both relations: $2 * (\text{Pages}_R + \text{Pages}_S)$
- * In match phase: Read both relations; $\text{Pages}_R + \text{Pages}_S$ I/Os
- Sort-Merge Join vs Hash Join:
- * Given a a reasonable amount of memory ($B > \sqrt{\text{Pages}_R}$), both have cost $3 * (\text{Pages}_R + \text{Pages}_s)$ I/Os
- * Hash join is superior if relation sizes differ greatly be-

- cause it needs less memory, and is also highly parallelizable
- * Sort-Merge is less sensitive to data skew and its result is sorted
- *More general join conditions:*
- * Equalities over several attributes: for INL, build index on composite key. For sort-merge and hash join, sort/hash-partition on the combination of the join columns
- * Inequality conditions: for INL, need (clustered!) B+ tree index; hash-join not usable; merge-join possible; block NL the best

Var	Desc
B	the number of data pages
R	number of records per page
D	average time to read or write a disk page
F	average fanout for a non-leaf page

	()	Scan	Equality	Range	Insert	Delete
Heap	BD	BD	$0.5BD$	BD	$2D$	$\text{Search} + D$
Sorted	BD	BD	$D \log_2 B$	$D(\log_2 B + \# \text{ matching pages})$	$\text{Search} + BD$	$\text{Search} + BD$
Clustered	$1.5BD$	$BD(R + 0.15)$	$D \log_F 1.5B$	$D(\log_F 1.5B + \# \text{ matching pages})$	$\text{Search} + D$	$\text{Search} + D$
Unclust. Tree	$BD(R + 0.15)$	$BD(R + 0.15)$	$D(1 + \log_F 0.15B)$	$D(\log_F 0.15B + \# \text{ matching pages})$	$\text{Search} + 2D$	$\text{Search} + 2D$
Unclust. Hash	$BD(R + 0.125)$	$BD(R + 0.125)$	$2D$	BD	$\text{Search} + 2D$	$\text{Search} + 2D$