

- **ACID**
 - **ATOMICITY:** An atomic transaction happens as one unit, either the whole thing commits or none of it does.
 - **CONSISTENCY:** A consistent transaction brings the DB from one valid state to another valid state with respect to any constraints.
 - **ISOLATION:** Concurrent isolated transactions would have the same result if run sequentially.
 - **DURABILITY:** A committed transaction will remain committed even in the event of a hardware failure.
- **RAID Levels**
 - Level 0: No redundancy (just stripin)
 - Level 1: Mirrored (two identical copies)
 - * Each disk has an exact mirror image
 - * Parallel reads; writes involve two disks
 - * Maximum transfer rate = transfer rate of one disk
 - Level 0+1 (Level 10): Striping and Mirroring
 - * Parallel reads; writes involve two disks
 - * Maximum transfer rate = aggregate bandwidth
 - Level 3: Bit-interleaved parity
 - * Striping Unit: one bit (or byte) (one check disk)
 - * Each read and write request involves all disks; disk array can process one request at a time
 - Level 4: Block-interleaved parity
 - * Striping unit: one disk block (one check disk)
 - * Parallel reads possible for small requests, large requests can utilize full bandwidth
 - * Writes involve modified block *and* check disk
 - Level 5: Block-interleaved distributed parity
 - * Similar to RAID level 4 but parity blocks are distributed over all disks
- **Buffer Management in a DBMS**
 - DBMS maintains buffer pool of frames, each frame holds a page, info is in **<frame#, pageid>** table
 - Choice of frame replacement dictated by replacement policy such as LRU
 - When a page is requested:
 - * If requested page is not in pool:
 - Choose a frame for replacement
 - If that frame is dirty, write it to disk
 - Read requested page into chosen frame
 - * Pin the page and return its address
 - * When done the requestor must indicate whether the page has been modified (dirty bit) and unpin
 - * Page in pool may be requested many times
 - A pin count is used and a page is a candidate for replacement iff **pin_count = 0**
 - Pinning increments pin count and unpinning decrements
 - * Concurrency control and recovery may entail additional I/O when a frame is chosen for replacement (write-ahead log protocol)
 - * Frame is chosen for replacement using LRU, clock, MRU, etc
 - * Sequential flooding: Caused by using LRU when the number of buffer frames is less than the number of pages in the file
 - **Files of Records**
 - Page or block is ok when doing I/O but higher levels of

- DBMS operate on *records* and thus want *files of records*
- **FILE:** A collection of pages each containing a collection of records. Must support
 - * Insert (append)/delete/modify record
 - * Read a particular record specified using *record id*
 - * Scan all records possibly with some conditions on the records to be retrieved
- **Unordered ‘‘Heap’’Files:**
 - * Simplest file structure that contains records in no particular (logical) order
 - * As file grows and shrinks, disk pages are allocated and de-allocated
 - * To support record-level operations we must:
 - Keep track of the *pages* in a file: **page id (pid)**
 - Keep track of the *free space* on a page
 - Keep track of the *records* on a page: **record id (rid)**
 - Keep track of *fields* within records
 - * Operations: create/destroy file, insert/delete record, fetch record with specific **rid**, scan all records
- Record formats: **Fixed Length**
 - * Information about field types is the same for all records in file; it is stored in *system catalogs*
 - * Finding the i^{th} field of a record does not require scanning the record
- Record formats: **Variable length**
 - * Several alternative formats (# of fields is fixed)
 - * Fields delimited by special symbols (e.g. \$ between fields)
 - * Fields preceded by lengths
- Record formats: **Variable length with directory**
 - * Use array of offsets at start of record
- **Heap file implemented as a list**
 - * The header page id and heap file name must be stored someplace
 - * Each page contains two extra pointers in this case
 - * Refinement: use several lists for different degrees of free space
- Page formats:
 - * File \sim_i collection of pages
 - * Page \sim_i collection of tuples/records
 - * Query operators deal with tuples
 - * Slotted page format:
 - Each page has a collection of *slots*
 - Each slot contains a record
 - * RID: **<page id, slot number>**
- **Heap file using a page directory**
 - * Page entries can include the number of free bytes on each page
 - * Directory is a collection of pages; linked list is one possible implementation
- **System catalogs:**
 - * For each relation:
 - name, file, file structure
 - name, type, length (if fixed) for each attribute
 - Index name, target, and kind for each index
 - also integrity constraints, defaults, nullability, etc
 - * For each index: structure (e.g. B+ tree) and search key fields

- * For each view: view name and definition (including query)
- * Plus statistics, authorization, buffer pool size, etc
- **Column Stores:**
 - * Store data “vertically”
 - * Contrast with a “row-store” that stores all the attributes of a tuple/record contiguously
 - * Each column can be stored as a separate file and compressed
 - * SAP HANA:
 - Dictionary compression per column
 - Column main: read-optimized store for immutable data. Uses high data compression and heuristic algorithms to order data to maximize secondary compression
 - Column delta: write-optimized store for inserts, updates, deletes. Uses less compression, appends updates to the end, and merges with main periodically.
 - * Additional types: prefix coding, run length coding, cluster coding, sparse coding, indirect coding
- **Indexes:**
 - * Speeds up selections on the search key fields for the index
 - * Contains a collection of data entries and supports efficient retrieval of all data entries k^* with a given key value k
- **B+ Tree Indexes**
 - * Leaf pages contain *data entries* and are chained (prev & next)
 - * Non-leaf pages have *index entries*, used to direct searches
 - * Insert/delete at $\log_F N$, keep tree *height-balanced* (F = fanout, N = # leaf pages)
 - * Minimum 50% occupancy (in all nodes except root). Each node contains $d \leq m \leq 2d$ entries; d = the *order* of the tree.
 - * Typical order $d = 100$
 - * Percentage of node that is full is more useful, typical fill-factor 67%
 - * Average *fanout* for non-leaves $F = 133$
 - * Inserting a data entry:
 - Find correct leaf L
 - Put data entry onto L
 - If L has enough space, done
 - Otherwise, must split L . Redistribute entries evenly, copy up the middle key (key must still exist in leaf). Insert index entry pointing to L_2 into parent of L .
 - This can happen recursively: if parent of L grows, need to push up middle key.
 - Splits “grow” the tree; root split increases height.
 - * Deleting a data entry:
 - Start at root, find leaf L where entry belongs
 - Remove the entry
 - If L is at least half full, done
 - Otherwise, if L has only $d - 1$ entries, try to redistribute, borrowing from sibling (adjacent node with same parent)
 - If redistribution fails, merge L and sibling
 - If merge occurred, must delete entry from parent

- (pointing to merged node)
- Merge can propagate to root, decreasing height of the tree

p	q	$p \text{ OR } q$	$p \text{ AND } q$	$p = q$
T	T	T	T	T
T	F	T	F	F
T	U	T	U	U
F	T	T	F	F
F	F	F	F	T
F	U	U	F	U
U	T	T	U	U
U	F	U	F	U
U	U	U	U	U

p	NOT p
T	F
F	T
U	U

Isolation Level		dirty reads		non-repeatable reads		phantoms	
SERIALIZABLE	READ UNCOMMITTED	Y	Y	Y	Y	Y	Y
	READ COMMITTED	N	N	Y	Y	Y	Y
	REPEATABLE READ	N	N	N	N	Y	Y
	Serializable	N	N	N	N	N	N